


```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: adult = pd.read_csv(r"C:\Users\786\Downloads\archive (5)\adult.csv")
adult
```

Out[2]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Blac
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	Whit
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	Whit
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Blac
4	18	?	103497	Some-college	10	Never-married	?	Own-child	Whit
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	Whit
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	Whit
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	Whit
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	Whit
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	Whit


48842 rows × 15 columns



```
In [3]: adult.head()
```

```
Out[3]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	ge
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Fe



```
In [4]: adult.describe()
```

```
Out[4]:
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.422382
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.391444
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000	99.000000

```
In [5]: adult["income"].describe()
```

```
Out[5]:
```

```
count      48842
unique         2
top      <=50K
freq      37155
Name: income, dtype: object
```

```
In [6]: adult.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
1   workclass             48842 non-null  object
2   fnlwgt               48842 non-null  int64
3   education            48842 non-null  object
4   educational-num      48842 non-null  int64
5   marital-status       48842 non-null  object
6   occupation           48842 non-null  object
7   relationship         48842 non-null  object
8   race                 48842 non-null  object
9   gender               48842 non-null  object
10  capital-gain         48842 non-null  int64
11  capital-loss         48842 non-null  int64
12  hours-per-week       48842 non-null  int64
13  native-country       48842 non-null  object
14  income               48842 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

```
In [7]: adult.nunique()
```

```
Out[7]: age                74
workclass                9
fnlwgt                 28523
education                16
educational-num         16
marital-status          7
occupation              15
relationship            6
race                   5
gender                 2
capital-gain           123
capital-loss           99
hours-per-week         96
native-country         42
income                 2
dtype: int64
```

```
In [8]: adult.index
```


```
Out[8]: RangeIndex(start=0, stop=48842, step=1)
```

```
In [9]: adult.columns
```

```
Out[9]: Index(['age', 'workclass', 'fnlwgt', 'education', 'educational-num',
               'marital-status', 'occupation', 'relationship', 'race', 'gender',
               'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
               'income'],
              dtype='object')
```


```
In [10]: sns.countplot(x="gender",data=adult)
```

```
Out[10]: <AxesSubplot:xlabel='gender', ylabel='count'>
```



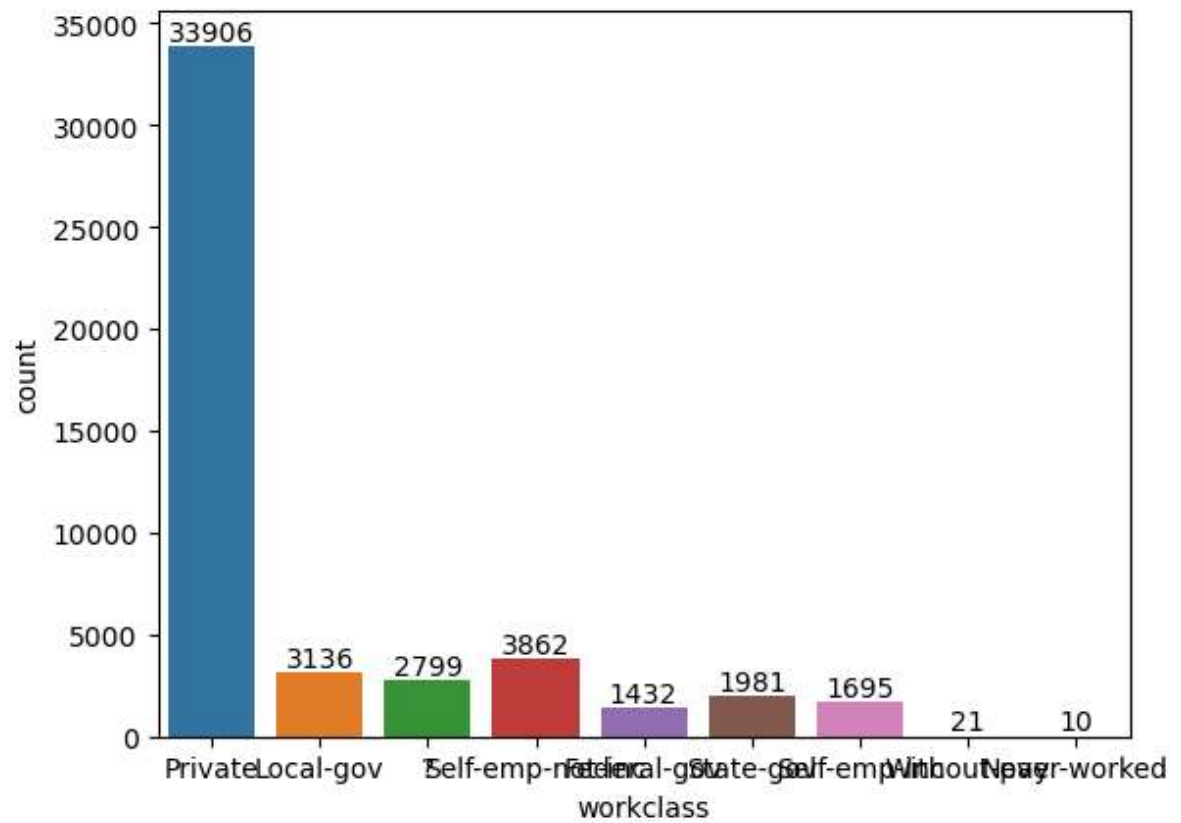
```
In [11]: ax = sns.countplot(x="gender", data=adult)
```

```
for bars in ax.containers:  
    ax.bar_label(bars)
```



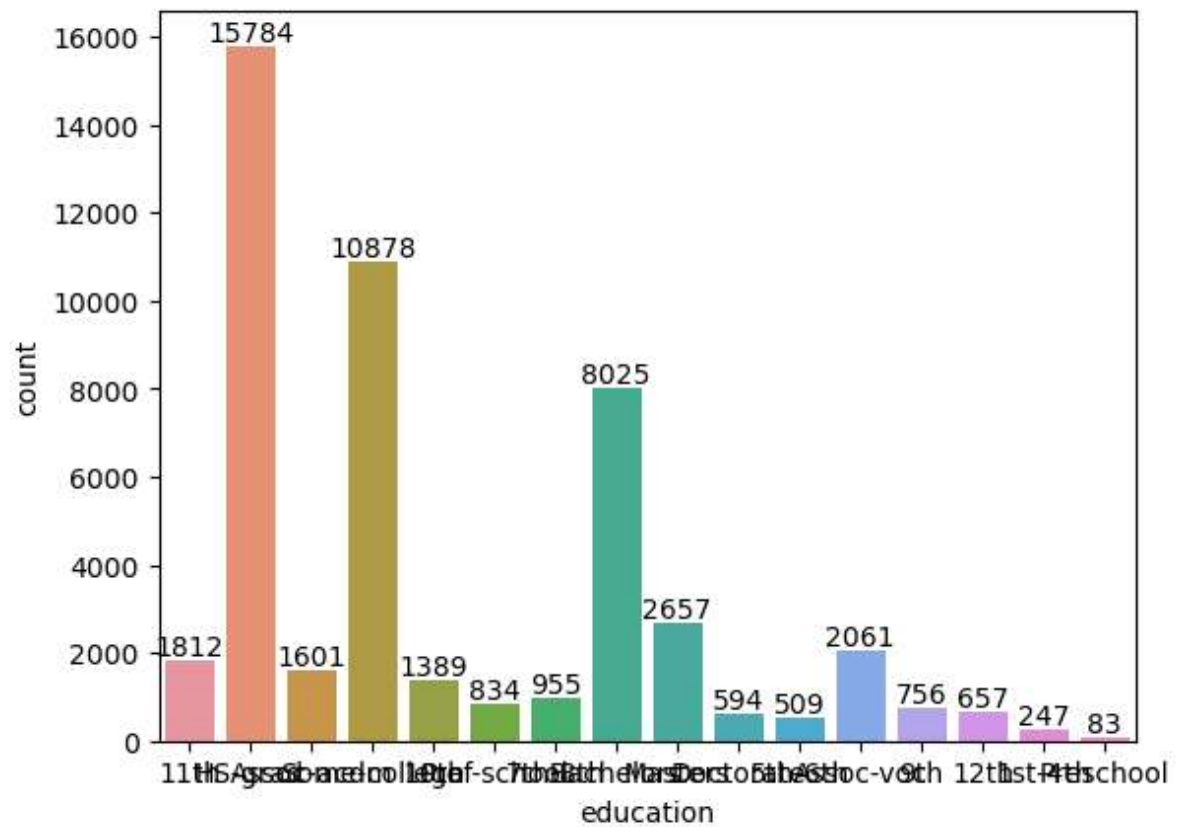
```
In [12]: ax = sns.countplot(x="workclass", data=adult)
```

```
for bars in ax.containers:  
    ax.bar_label(bars)
```



```
In [13]: ax = sns.countplot(x="education", data=adult)
```

```
for bars in ax.containers:  
    ax.bar_label(bars)
```



In [14]: adult

Out[14]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	rac
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Blac
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	Whit
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	Whit
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Blac
4	18	?	103497	Some-college	10	Never-married	?	Own-child	Whit
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	Whit
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	Whit
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	Whit
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	Whit
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	Whit

48842 rows × 15 columns




```
In [15]: adult.isnull()
```

```
Out[15]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race
0	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False
...
48837	False	False	False	False	False	False	False	False	False
48838	False	False	False	False	False	False	False	False	False
48839	False	False	False	False	False	False	False	False	False
48840	False	False	False	False	False	False	False	False	False
48841	False	False	False	False	False	False	False	False	False

48842 rows × 15 columns



```
In [16]: adult.isnull().sum()
```

```
Out[16]: age                0
workclass                0
fnlwgt                  0
education                0
educational-num         0
marital-status          0
occupation              0
relationship            0
race                   0
gender                  0
capital-gain            0
capital-loss            0
hours-per-week          0
native-country          0
income                 0
dtype: int64
```

```
In [17]: len(adult)
```

```
Out[17]: 48842
```

```
In [18]: adult["marital-status"].describe()
```

```
Out[18]: count          48842
unique              7
top      Married-civ-spouse
freq          22379
Name: marital-status, dtype: object
```


```
In [19]: adult["marital-status"].unique()
```

```
Out[19]: array(['Never-married', 'Married-civ-spouse', 'Widowed', 'Divorced',
        'Separated', 'Married-spouse-absent', 'Married-AF-spouse'],
        dtype=object)
```

```
In [20]: adult.head()
```

```
Out[20]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	ge
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Fe



In [21]: adult

Out[21]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	rac
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Blac
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	Whit
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	Whit
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Blac
4	18	?	103497	Some-college	10	Never-married	?	Own-child	Whit
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	Whit
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	Whit
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	Whit
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	Whit
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	Whit

48842 rows × 15 columns



```
In [22]: adult.head()
```

```
Out[22]:
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female

```
In [23]: adult=adult.drop(["education", "fnlwgt", "marital-status"],axis=1)
adult.head()
```

```
Out[23]:
```

	age	workclass	educational-num	occupation	relationship	race	gender	capital-gain	capital-loss	hours per week
0	25	Private	7	Machine-op-inspct	Own-child	Black	Male	0	0	41
1	38	Private	9	Farming-fishing	Husband	White	Male	0	0	51
2	28	Local-gov	12	Protective-serv	Husband	White	Male	0	0	41
3	44	Private	10	Machine-op-inspct	Husband	Black	Male	7688	0	41
4	18	?	10	?	Own-child	White	Female	0	0	31

```
In [24]: adult.replace('?', np.NaN,inplace = True)
adult.head()
```

Out[24]:

	age	workclass	educational- num	occupation	relationship	race	gender	capital- gain	capital- loss	hours per week
0	25	Private	7	Machine- op-inspct	Own-child	Black	Male	0	0	41
1	38	Private	9	Farming- fishing	Husband	White	Male	0	0	51
2	28	Local-gov	12	Protective- serv	Husband	White	Male	0	0	41
3	44	Private	10	Machine- op-inspct	Husband	Black	Male	7688	0	41
4	18	NaN	10	NaN	Own-child	White	Female	0	0	31

```
In [25]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
```

```
In [26]: adult["workclass"] = le.fit_transform(adult["workclass"])
adult["occupation"] = le.fit_transform(adult["occupation"])
adult["relationship"] = le.fit_transform(adult["relationship"])
adult["race"] = le.fit_transform(adult["race"])
adult["gender"] = le.fit_transform(adult["gender"])
adult["native-country"] = le.fit_transform(adult["native-country"])
adult["income"] = le.fit_transform(adult["income"])


adult.head()
```

Out[26]:

	age	workclass	educational- num	occupation	relationship	race	gender	capital- gain	capital- loss	hours- per- week
0	25	3	7	6	3	2	1	0	0	40
1	38	3	9	4	0	4	1	0	0	50
2	28	1	12	10	0	4	1	0	0	40
3	44	3	10	6	0	2	1	7688	0	40
4	18	8	10	14	3	4	0	0	0	30


```
In [27]: sns.barplot(x = 'income', y='age', data=adult)
```

```
Out[27]: <AxesSubplot:xlabel='income', ylabel='age'>
```



```
In [28]: sns.heatmap(adult.corr())
```

```
Out[28]: <AxesSubplot:>
```



```
In [29]: x = adult.drop(['income'],axis=1)
y = adult['income']
```

```
In [30]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [31]: from sklearn.ensemble import RandomForestClassifier
Clf = RandomForestClassifier()
```

```
In [32]: Clf.fit(x_train,y_train)
```

```
Out[32]: RandomForestClassifier()
```

```
In [33]: y_pred = Clf.predict(x_test)
y_pred
```

```
Out[33]: array([0, 0, 1, ..., 0, 0, 0])
```

```
In [34]: !pip install scikit-learn
```

```
Defaulting to user installation because normal site-packages is not writeable
Requirement already satisfied: scikit-learn in c:\programdata\anaconda3\lib\site-packages (1.0.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn) (2.2.0)
Requirement already satisfied: scipy>=1.1.0 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn) (1.9.1)
Requirement already satisfied: joblib>=0.11 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn) (1.1.0)
Requirement already satisfied: numpy>=1.14.6 in c:\programdata\anaconda3\lib\site-packages (from scikit-learn) (1.21.5)
```

```
In [35]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.89	0.92	0.91	11236
1	0.71	0.63	0.67	3417
accuracy			0.85	14653
macro avg	0.80	0.78	0.79	14653
weighted avg	0.85	0.85	0.85	14653

```
In [36]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,y_pred)
```

```
Out[36]: array([[10360,   876],
               [ 1269,  2148]], dtype=int64)
```

```
In [37]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred)
```

```
Out[37]: 0.8536135944857708
```

```
In [38]: x = adult.drop(['income'],axis=1)
y = adult['income']
```

```
In [39]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [40]: from sklearn.naive_bayes import GaussianNB
gb = GaussianNB()
```

```
In [41]: gb.fit(x_train,y_train)
```

```
Out[41]: GaussianNB()
```



```
In [42]: y_pred = gb.predict(x_test)
y_pred
```

```
Out[42]: array([0, 1, 1, ..., 0, 0, 0])
```

```
In [43]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,y_pred)
```

```
Out[43]: array([[10613,   557],
               [ 2405,  1078]], dtype=int64)
```

```
In [44]: from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred)
```

```
Out[44]: 0.797857094110421
```

```
In [45]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.82	0.95	0.88	11170
1	0.66	0.31	0.42	3483
accuracy			0.80	14653
macro avg	0.74	0.63	0.65	14653
weighted avg	0.78	0.80	0.77	14653

```
In [ ]:
```